

ANaConDA: A Framework for Analysing Multi-threaded C/C++ Programs on the Binary Level

Jan Fiedor, Tomáš Vojnar

Brno University of Technology (BUT)

RV, September 26, 2012

- Why to analyse multi-threaded C/C++ programs?
 - There is an ever-growing number of such programs
 - Concurrency errors easy to create, but difficult to discover
 - Many tools for Java (IBM ConTest, RoadRunner, CalFuzzer, Chord, etc.), however, not so many for C/C++ (Fjalar)
- Dynamic analysis
 - Extrapolates the witnessed behaviour
 - May detect errors not witnessed in the given execution
 - Needs to monitor the execution of a program (difficult)
- ANaConDA framework
 - A framework simplifying the creation of dynamic analysers
 - Monitors the execution of a multi-threaded C/C++ program
 - Offers notification about important events to the analyser
 - Supports noise injection

Information Provided by the Framework

Memory access information:

- Reads, writes and atomic updates

Synchronisation information:

- Lock acquisitions/releases, signaling conditions and waiting on them

Thread information:

- Thread started/finished

Exception information:

- Exception thrown/caught

Convenient information for localising detected errors:

- Backtraces (containing return addresses)

A Simple Analyser Monitoring Lock Operations

```
PLUGIN_INIT_FUNCTION()
{
    // Register a callback function called before a lock is released
    SYNC_BeforeLockRelease(beforeLockRelease);

    // Register a callback function called after a lock is acquired
    SYNC_AfterLockAcquire(afterLockAcquire);
}

VOID beforeLockRelease(THREADID tid, LOCK lock)
{
    CONSOLE("Before lock released: thread " + decstr(tid) + ", lock " + lock + "\n");
}

VOID afterLockAcquire(THREADID tid, LOCK lock)
{
    CONSOLE("After lock acquired: thread " + decstr(tid) + ", lock " + lock + "\n");
}
```

Instantiation for a Specific Multithreading Library

The framework might be **configured** to work with **various** multithreading libraries:

- The same dynamic analyser may be **reused** for different libraries
- Currently instantiated for the **pthread** library and **Win32 API**

The **user** must specify:

- Which **functions** perform certain types of thread-related operations
 - Specify the **names** of these functions
- Which arguments represent the **synchronisation resources**
 - Specify the **indices** of these arguments
- How to **transform** synchronisation resources to their abstract identifications
 - Specify the **mapper objects** used to perform the transformation

Instantiation for the Pthread Library

To configure monitoring of **lock acquisitions**:

- The required information must be added to `conf/hooks/lock`

```
# <function name>      <index>  <mapper object>
pthread_mutex_lock      1          addr()
pthread_mutex_trylock   1          addr()
```

To configure monitoring of **lock releases**:

- The required information must be added to `conf/hooks/unlock`

```
# <function name>      <index>  <mapper object>
__pthread_mutex_unlock_usercnt  1          addr()
```

Noise Injection Techniques:

- Aim at **increasing** the number of different witnessed **interleavings**
- Disturb the **scheduling** of threads by inserting noise generating code
 - e.g., by inserting calls of `yield` or `sleep`
- Force the program to **switch threads** at times it would normally seldom do it

User may typically influence:

- **Type** of noise (currently *sleep* and *yield* noise is supported)
- Noise **frequency** (how often the noise should occur)
- Noise **strength** (how strong the noise should be)

Configuring Noise Injection

The framework supports **fine-grained** combinations of noise, i.e., **different** noise injection settings might be used for:

- **Reads, writes** and **atomic updates** (configured in `anaconda.conf`)
- Each of the **monitored functions** (configured in `lock, unlock` etc.)

An example how to use different noise settings for **reads** and **writes**:

```
[noise]                # Global noise settings
type = yield           # Insert calls to yield
frequency = 100         # Inject noise in 10 % of times
strength = 4           # Give up the CPU 4 times
[noise.read]           # Noise settings for read accesses
type = yield           # Insert calls to yield
frequency = 200        # Inject noise in 20 % of times
strength = 8           # Give up the CPU 8 times
[noise.write]          # Noise settings for write accesses
type = sleep           # Insert calls to sleep
frequency = 400        # Inject noise in 40 % of times
strength = 2           # Sleep for 2 milliseconds
```


Experiments

ANaConDA is a *pintool* (plugin for the PIN framework):

- Instrumentation done in the **memory** (transparent use of **libraries**)
- Can handle **generated** and **self-modifying** code
- May be used on both **Linux** and **Windows**
- **Slowdown** of the execution is around 100 times

Firefox 10 browser

- So far without a test harness
- Found several **known** data races considered as **harmless**
- Proved that the tool can handle even **very large** programs

Unicap libraries: libraries for concurrent video processing

- Found several **previously unknown** data races
- Some of them cause programs using these libraries to **crash**

- We have presented ANaConDA
 - A framework simplifying the creation of dynamic analysers
- We have shown
 - How to write an analyser
 - How to instantiate the framework for a particular multithreading library
 - How to configure and use the noise injection

- Improvements to the [framework](#)
 - Additional [notifications](#) (fork/join notifications etc.)
 - Better access to [debugging information](#)
 - Smarter [instrumentation](#)
 - ...
- More experiments
- More [sophisticated](#) types of noises
- [New detectors](#) for concurrency errors

Thank you for your attention!

ANaConDA framework available at:

<http://www.fit.vutbr.cz/research/groups/verifit/tools/anaconda/>

An Example How to Obtain a Backtrace

```
VOID beforeMemoryWrite(THREADID tid, ADDRINT addr, UINT32 size,
    const VARIABLE& variable, const LOCATION& location)
{
    // Helper variables
    Backtrace bt;
    Symbols symbols;

    // Get the backtrace of the current thread
    THREAD_GetBacktrace(tid, bt);
    // Translate the return addresses to locations
    THREAD_GetBacktraceSymbols(bt, symbols);

    CONSOLE_NOPREFIX("Thread " + decstr(tid) + " backtrace:\n");

    for (Symbols::size_type i = 0; i < symbols.size(); i++)
    { // Print information about each return address in the backtrace
        CONSOLE_NOPREFIX("    #" + decstr(i) + (i > 10 ? " " : " ")
            + symbols[i] + "\n");
    }
}
```

Choosing Parts of a Program To Be Monitored

Any image (executable file, library) can be:

- **Excluded** from instrumentation (conf/filters/ins/exclude)
- Forced to be **instrumented** (conf/filters/ins/include)

An example how to configure the framework to **not** instrument **libraries**:

```
# Do not instrument standard Linux libraries
/lib/*
/lib64/*
/usr/lib/*
/usr/local/lib/*
# Do not instrument standard Windows libraries
${windir}/system32/*
${windir}/SYSTEM32/*
# Do not instrument PIN libraries
${PIN_HOME}/*
# Do not instrument ANaConDA plugins
*/lib/ia32/*
*/lib/intel64/*
```

Performing an Analysis

To **analyse** a multi-threaded C/C++ program using the **framework**, one can use the `anaconda.sh` (`anaconda.bat`) script:

```
./anaconda.[sh|bat] <path-to-analyser> <path-to-program>
```

To pass **additional parameters** to PIN or ANaConDA, one has to use the following command:

```
$PIN_HOME/pin[.bat] <pin-args> -t <path-to-anaconda>  
-a <path-to-analyser> <anaconda-args>  
-- <path-to-program> <program-args>
```